Evidence Verification Ontology — Sample

Tahir Tidjani · Anonymized extract from a stealth AI-verification product · Shared vocabulary that lets AI agents, schemas and human reviewers mean the same thing by “claim”, “evidence” and “source”.

CORE ENTITY TYPES (17 OF 26 SHOWN) AND AUDIT RECORD

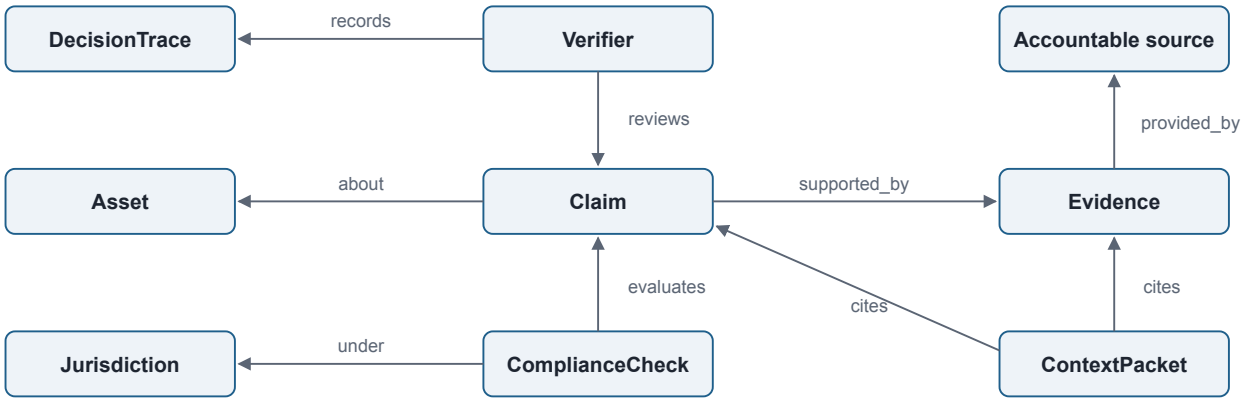
Parties
`Issuer`
`Operator`
`Supplier`
`Custodian`
`Verifier`

Subject
`Asset`
`Jurisdiction`
`License`
`Credential`

Evidence & assessment
`Claim`
`Evidence`
`ComplianceCheck`
`RiskScore`
`Valuation`
`ProductionEvent`
`Dispute`

AI context & audit
`ContextPacket`
`DecisionTrace`

EVIDENCE GRAPH (DESIGNED RELATIONSHIPS)



PROVENANCE & CONTROL FIELDS

`asset_id · claim_id evidence_id`	Stable identifiers used identically across agents, JSON schemas and data models.
`source_evidence source_claims`	Lineage: every conclusion keeps a pointer to what it was derived from.
`confidence_score`	Model confidence, explicitly *not* legal certainty.
`requires_human_review`	Mandatory flag on high-risk AI outputs; a person decides.
`review_status`	Workflow state: accepted, manual review, rejected.

DESIGN RULES

- **AI assesses, people attest.** The vocabulary forbids wording such as “AI-certified” or “sensor-proven”.
- **Truth anchor:** each fact traces to the accountable source closest to it.
- **Smallest reliable context:** agents receive a ContextPacket of cited claims, not whole repositories.
- **Decision trace:** retrieved context, tool calls, policy checks, citations and approvals are recorded for audit.
- **One name, one meaning:** extend the glossary rather than invent new terms in code.

Status. Entity names and field conventions are used in the prototype’s agent prompts, JSON schemas and data models. The relationships above are the evidence-graph design and are not yet implemented as a graph store. Finance-specific entity types are omitted from this extract.